

# Approaches to making R faster

Philip Barrett

University of Chicago

2 May 2014

# Installation

You should have installed everything by now.

If not, do so while I'm rambling on (next  $\sim$  5 minutes).

Extra: you may want to install microbenchmark:

```
install.packages('microbenchmark')  
library(microbenchmark)
```

# The limitations of R

- ▶ R is great in many ways: flexible, easy to write, mature.
- ▶ But it isn't the fastest.
- ▶ So what if your problem has large computational demands?
- ▶ Two approaches:
  1. Parallelize
  2. Integrate other languages
- ▶ Today, focus on 2 (esp. C++), but discuss 1 briefly.

# Parallelism

- ▶ Pros:
  - ▶ Easy: implementation and learning
  - ▶ Time cost reduced by factor of  $\sim 50$  on small(ish) systems
  - ▶ No (software) cost of extra nodes on v large systems
- ▶ Cons:
  - ▶ At best,  $t \propto 1/N$
  - ▶ Communication time
  - ▶ Sequential operations
- ▶ Resources:
  - ▶ “Parallel R”, McCallum & Weston (O'Reilly)
  - ▶ Steven Mohr: [smohr@uchicago.edu](mailto:smohr@uchicago.edu)

# C++

- ▶ Pros:
  - ▶ Fast (esp. for loops & Linear algebra)
  - ▶ Universal: The “Latin” of CS
  - ▶ Mature
- ▶ Cons:
  - ▶ Harder to learn
  - ▶ Slower to write
  - ▶ Compilers
- ▶ Our approach:
  - ▶ Use the Rcpp and RcppArmadillo packages
  - ▶ Integrates beautifully to R.
  - ▶ Call functions in R, but written in C++.
  - ▶ No painful set up or compilation headaches.

Why is C++ so fast?

## Why is C++ so fast?

*Blah blah blah blah blah blah blah blah  
blah blah **defining types** blah blah blah blah  
blah blah blah blah blah blah blah blah  
blah blah blah blah blah blah blah blah  
blah blah blah **machine code** blah blah blah  
blah blah blah blah blah blah blah blah  
blah **compilers** blah blah blah blah blah  
blah blah blah blah blah blah blah blah  
blah blah blah blah blah blah blah blah  
blah blah **but basically, it just is.***

# The motherlode

Combining C++ and parallelism can give you *massive* speed gains.

- ▶ Strategy:
  1. Write R code that parallelizes a repeated operation
  2. Write the code for that operation in C++
- ▶ Conservative numbers:
  - ▶ Parallelism gives a  $\times 50$  speed gain
  - ▶ C++ operation is  $\sim 100$  times faster
  - ▶ Total speed gain  $\sim 5000$
- ▶ Excluding cluster initialization, this means:
  1. Before: 1 hour, after:  $\sim 0.7$  seconds
  2. Before: 1 day, after:  $\sim 17$  seconds
  3. Before: 1 week, after:  $\sim 2$  mins
  4. Before: 1 year, after:  $\sim 1.5$  hours

In practice, speed gains from C++ part be much greater.

NB: Parallelizing in C++ alone is not fun.



# Outline

- ▶ Example
- ▶ Basics of coding in C++
  1. Hello world
  2. Types
  3. Returns
  4. If..else
  5. Loops
  6. R-style vectors (with zero-based counting!)
  7. R-style matrices and lists
- ▶ Matrix algebra with RcppArmadillo.
  1. Vectors and matrices
  2. Simple operations: Addition, multiplication, matrix metadata
  3. Harder operations: Accessing submatrices, inversion
- ▶ Rewrite John Eric's MCMC example