

Approaches to making R faster

Philip Barrett

University of Chicago

2 May 2014

Installation

You should have installed everything by now.

If not, do so while I'm rambling on (next \sim 5 minutes).

The limitations of R

- ▶ R is great in many ways: flexible, easy to write, mature.
- ▶ But it isn't the fastest.
- ▶ So what if your problem has large computational demands?
- ▶ Two approaches:
 1. Parallelize
 2. Integrate other languages
- ▶ Today, focus on 2 (esp. C++), but discuss 1 briefly.

Parallelism

- ▶ Pros:
 - ▶ Easy: implementation and learning
 - ▶ Time cost reduced by factor of ~ 50 on small(ish) systems
 - ▶ No (software) cost of extra nodes on v large systems
- ▶ Cons:
 - ▶ At best, $t \propto 1/N$
 - ▶ Communication time
 - ▶ Sequential operations
- ▶ Resources:
 - ▶ “Parallel R”, McCallum & Weston (O'Reilly)
 - ▶ Steven Mohr: smohr@uchicago.edu

C++

- ▶ Pros:
 - ▶ Fast (esp. for loops & Linear algebra)
 - ▶ Universal: The “Latin” of CS
 - ▶ Mature
- ▶ Cons:
 - ▶ Harder to learn
 - ▶ Slower to write
 - ▶ Compilers
- ▶ Our approach:
 - ▶ Use the `Rcpp` and `RcppArmadillo` packages
 - ▶ Integrates beautifully to R .
 - ▶ Call functions in R, but written in C++.
 - ▶ No painful set up or compilation headaches.

Why is C++ so fast?

Why is C++ so fast?

*Blah blah blah blah blah blah blah blah
blah blah **defining types** blah blah blah blah
blah blah blah blah blah blah blah blah
blah blah blah blah blah blah blah blah
blah blah blah **machine code** blah blah blah
blah blah blah blah blah blah blah blah
blah **compilers** blah blah blah blah blah
blah blah blah blah blah blah blah blah
blah blah blah blah blah blah blah blah
blah blah **but basically, it just is.***

The motherlode

Combining C++ and parallelism can give you *massive* speed gains.

- ▶ Strategy:
 1. Write R code that parallelizes a repeated operation
 2. Write the code for that operation in C++
- ▶ Conservative numbers:
 - ▶ Parallelism gives a $\times 50$ speed gain
 - ▶ C++ operation is ~ 60 times faster
 - ▶ Total speed gain ~ 3000
- ▶ Excluding cluster initialization, this means:
 1. Before: 1 hour, after: ~ 0.02 seconds
 2. Before: 1 day, after: ~ 30 seconds
 3. Before: 1 week, after: ~ 3 mins
 4. Before: 1 year, after: ~ 3 hours

In practice, speed gains from C++ part be much greater.

NB: Parallelizing in C++ alone is not fun.

Outline

- ▶ Example
- ▶ Basics of coding in C++
 1. Hello world
 2. Types
 3. Returns
 4. If..else
 5. Loops
 6. R-style vectors (with zero-based counting!)
 7. R-style matrices and lists
- ▶ Matrix algebra with Rcpp Armadillo .
 1. Vectors and matrices
 2. Simple operations: Addition, multiplication, matrix metadata
 3. Harder operations: Accessing submatrices, inversion
- ▶ Rewrite John Eric's MCMC example